

SOMESH ZANWAR

Product & Data Analyst • SQL • A/B Experimentation • KPI & Funnel Analytics • AI/Agent Governance

someshzanwar345@gmail.com | (682) 403-0698 | [LinkedIn](#) | [Portfolio](#) | [GitHub](#) | Arlington, TX | F-1 OPT (CPT Available)

PROFESSIONAL SUMMARY

Data & Product Analyst with 1+ year of industry experience building analytics pipelines, A/B experimentation frameworks, and AI/agent governance systems. Delivered 30% improvement in reporting accuracy and 40% reduction in data preparation time in production environments. Currently completing M.S. in Data Science at UT Arlington (GPA 3.8); active contributor to Microsoft's Agent Governance Toolkit open-source project. Authorized to work in the US on F-1 OPT: no sponsorship required for 3 years.

TECHNICAL SKILLS

Analytics & Querying: SQL (Window Functions, CTEs, Query Optimization, Joins), Python (Pandas, NumPy), BigQuery

Product & Business Analytics: KPI Development, Funnel Analysis, Cohort Analysis, Retention Analysis, A/B Testing, Experiment Design

Visualization & BI: Tableau, Power BI, Streamlit, Plotly: Dashboard Design, Stakeholder Reporting, Data Storytelling

Data Engineering: ETL Pipelines, dbt, Data Modeling (Star/Snowflake Schema), PostgreSQL, MySQL, Data Quality Monitoring, Agent Governance (Microsoft AGT)

Machine Learning (Applied): Scikit-learn, XGBoost, Feature Engineering: supporting product experimentation & prediction workflows

Cloud & Tools: Google Cloud Platform, Docker, Git, Jira, Salesforce

WORK EXPERIENCE

Junior Data Analyst — IVS Software Solutions

Jun 2023 – Sep 2024

Pune, India

- Analyzed 80K–150K operational records per cycle using SQL and Python to surface KPI trends, detect anomalies, and power stakeholder reporting, directly reducing decision latency for 3 internal business teams.
- Conducted structured exploratory data analysis (EDA) that uncovered systematic data quality gaps, improving downstream reporting accuracy by 30% and establishing recurring data validation checkpoints.
- Re-engineered SQL pipelines with optimized join logic, window functions, and incremental loads, cutting data preparation time by 40% and accelerating dashboard refresh cycles from daily to near-real-time.
- Built reusable, analysis-ready datasets supporting weekly and monthly reporting workflows, eliminating 4–6 hours of manual data wrangling per stakeholder cycle.

PROJECTS

Data Quality-Aware Agent Governance | Python, Microsoft AGT, PostgreSQL, dbt

- Built a two-layer governance prototype using Microsoft's Agent Governance Toolkit, enforcing both agent authorization and dataset trustworthiness checks, blocking agent actions when target datasets fail freshness or quality thresholds, not only when the agent lacks permission.
- Merged contributor to Microsoft's Agent Governance Toolkit: PR #1818 accepted by principal maintainer; project listed in official ADOPTERS.md alongside Microsoft internal teams.

AI Data Governance & Quality Monitoring System | PostgreSQL, Python, dbt, Power BI

- Engineered a governance layer monitoring 150K+ daily events across ingestion pipelines, enforcing schema validation and automated cross-source reconciliation for 90% of data sources.
- Deployed real-time anomaly detection cutting incident response time by 25%; built AI-powered Power BI dashboard translating data failures into actionable explanations for non-technical stakeholders.

Decision Intelligence Experimentation Platform | PostgreSQL, Python, Streamlit

- Architected an end-to-end A/B testing platform simulating 50K+ users and 280K+ events; designed hypothesis testing pipelines (z-test, t-test, Mann-Whitney) surfacing a conversion-vs-revenue trade-off that informed a go/no-go product launch decision.
- Designed dbt-powered metrics layer and Streamlit self-serve dashboard, reducing analyst involvement in experiment readouts by ~60%.

Metric Decomposition Engine | PostgreSQL, Python, Pandas, Streamlit

- Designed an automated root-cause analysis system decomposing event-based metrics across dimensions on 1.3M+ GitHub Archive events, ranking each segment's contribution to total metric movement.
- Generated plain-English incident reports via Jinja2 and an interactive Streamlit dashboard, turning multi-hour manual metric investigations into a self-serve workflow running in seconds.

EDUCATION

Masters in Data Science

Jan 2025 – Dec 2026

The University of Texas at Arlington, TX | GPA: 3.8

Bachelor of Engineering in Artificial Intelligence and Machine Learning

Aug 2020 – May 2024

Savitribai Phule Pune University, India | GPA: 8.2