

SOMESH ZANWAR

Product & Data Analyst • SQL • A/B Experimentation • KPI & Funnel Analytics • AI/Agent Governance
someshzanwar345@gmail.com | (682) 403-0698 | [LinkedIn](#) | [Portfolio](#) | [GitHub](#) | Arlington, TX | F-1 OPT (CPT Available)

PROFESSIONAL SUMMARY

Data & Product Analyst with 1+ year of experience using SQL, Python, and BI tools to build reporting workflows, analyze operational data, and improve data quality. Built governance-focused analytics projects using PostgreSQL, dbt, Power BI, and Microsoft Agent Governance Toolkit. M.S. Data Science candidate at UT Arlington with three merged AGT contributions across adopter documentation, tutorial verification, and a data-quality-aware governance example. Authorized to work in the U.S. on F-1 OPT; CPT available.

TECHNICAL SKILLS

Analytics & Querying: SQL (Window Functions, CTEs, Query Optimization, Joins), Python (Pandas, NumPy), BigQuery

Product & Business Analytics: KPI Development, Funnel Analysis, Cohort Analysis, Retention Analysis, A/B Testing, Experiment Design

Visualization & BI: Tableau, Power BI, Streamlit, Plotly: Dashboard Design, Stakeholder Reporting, Data Storytelling

Data Engineering: ETL Pipelines, dbt, Data Modeling (Star/Snowflake Schema), PostgreSQL, MySQL, Data Quality Monitoring, Agent Governance (Microsoft AGT)

Machine Learning (Applied): Scikit-learn, XGBoost, Feature Engineering: supporting product experimentation & prediction workflows

Cloud & Tools: Google Cloud Platform, Docker, Git, Jira, Salesforce

WORK EXPERIENCE

Junior Data Analyst — IVS Software Solutions Pune, India

Jun 2023 – Sep 2024

- Analyzed 80K–150K operational records per cycle using SQL and Python to surface KPI trends, detect anomalies, and power stakeholder reporting, directly reducing decision latency for 3 internal business teams.
- Conducted structured exploratory data analysis (EDA) that uncovered systematic data quality gaps, improving downstream reporting accuracy by 30% and establishing recurring data validation checkpoints.
- Re-engineered SQL pipelines with optimized join logic, window functions, and incremental loads, cutting data preparation time by 40% and accelerating dashboard refresh cycles from daily to near-real-time.
- Built reusable, analysis-ready datasets supporting weekly and monthly reporting workflows, eliminating 4–6 hours of manual data wrangling per stakeholder cycle.

PROJECTS

Data Quality-Aware Agent Governance | Python, Microsoft AGT, PostgreSQL, dbt

- Built a two-layer governance prototype using Microsoft AGT that combines agent authorization with dataset trust checks, blocking actions when target data fails freshness or quality thresholds.
- Created audit-ready decision logic showing whether an action was blocked by agent policy or data quality checks.

AI Data Governance & Quality Monitoring System | PostgreSQL, Python, dbt, Power BI

- Engineered a governance layer monitoring 150K+ daily events across ingestion pipelines, enforcing schema validation and systematized cross-source reconciliation for 90% of data sources.
- Deployed real-time anomaly detection cutting incident response time by 25%; built AI-powered Power BI dashboard translating data failures into actionable explanations for non-technical stakeholders.

Decision Intelligence Experimentation Platform | PostgreSQL, Python, Streamlit

- Architected an end-to-end A/B testing platform simulating 50K+ users and 280K+ events; designed hypothesis testing pipelines (z-test, t-test, Mann-Whitney) surfacing a conversion-vs-revenue trade-off that informed a go/no-go product launch decision.
- Designed dbt-powered metrics layer and Streamlit self-serve dashboard, reducing analyst involvement in experiment readouts by ~60%.

Metric Decomposition Engine | PostgreSQL, Python, Pandas, Streamlit

- Developed an automated root-cause analysis system decomposing event-based metrics across dimensions on 1.3M+ GitHub Archive events, ranking each segment's contribution to total metric movement.
- Generated plain-English incident reports via Jinja2 and an interactive Streamlit dashboard, turning multi-hour manual metric investigations into a self-serve workflow running in seconds.

EDUCATION

Masters in Data Science

Jan 2025 – Dec 2026

The University of Texas at Arlington, TX | **GPA: 3.67**

Bachelor of Engineering in Artificial Intelligence and Machine Learning

Aug 2020 – May 2024

Savitribai Phule Pune University, India | **GPA: 8.2**

OPEN-SOURCE CONTRIBUTIONS

Microsoft Agent Governance Toolkit | Open Source Contributor

- Contributed three merged PRs to Microsoft AGT: #1818 added Data Quality-Aware Agent Governance to ADOPTERS.md, #2038 fixed Agent SRE tutorial issues after Windows verification, and #2224 added a data-quality-aware governance example under examples/.
- Implemented a two-layer governance pattern combining AGT policy evaluation with dataset trust signals, blocking agent actions when target data is stale, low quality, or failing validation.