

SOMESH ZANWAR

Product & Data Analyst • SQL • A/B Experimentation • KPI & Funnel Analytics • AI/Agent Governance

someshzanwar345@gmail.com | (682) 403-0698 | [LinkedIn](#) | [Portfolio](#) | [GitHub](#) | Arlington, TX | F-1 OPT (CPT Available)

PROFESSIONAL SUMMARY

Data & Product Analyst with 1+ year of experience using SQL, Python, and BI tools to build reporting workflows, analyze operational data, and improve data quality. Built governance-focused analytics projects using PostgreSQL, dbt, Power BI, and Microsoft Agent Governance Toolkit. M.S. Data Science candidate at UT Arlington with merged open-source contributions to Microsoft's Agent Governance Toolkit, including adopter documentation and tutorial fixes from end-to-end verification. Authorized to work in the U.S. on F-1 OPT; CPT available.

TECHNICAL SKILLS

Analytics & Querying: SQL (Window Functions, CTEs, Query Optimization, Joins), Python (Pandas, NumPy), BigQuery

Product & Business Analytics: KPI Development, Funnel Analysis, Cohort Analysis, Retention Analysis, A/B Testing, Experiment Design

Visualization & BI: Tableau, Power BI, Streamlit, Plotly: Dashboard Design, Stakeholder Reporting, Data Storytelling

Data Engineering: ETL Pipelines, dbt, Data Modeling (Star/Snowflake Schema), PostgreSQL, MySQL, Data Quality Monitoring, Agent Governance (Microsoft AGT)

Machine Learning (Applied): Scikit-learn, XGBoost, Feature Engineering: supporting product experimentation & prediction workflows

Cloud & Tools: Google Cloud Platform, Docker, Git, Jira, Salesforce

WORK EXPERIENCE

Junior Data Analyst — IVS Software Solutions

Jun 2023 – Sep 2024

Pune, India

- Analyzed 80K–150K operational records per cycle using SQL and Python to surface KPI trends, detect anomalies, and power stakeholder reporting, directly reducing decision latency for 3 internal business teams.
- Conducted structured exploratory data analysis (EDA) that uncovered systematic data quality gaps, improving downstream reporting accuracy by 30% and establishing recurring data validation checkpoints.
- Re-engineered SQL pipelines with optimized join logic, window functions, and incremental loads, cutting data preparation time by 40% and accelerating dashboard refresh cycles from daily to near-real-time.
- Built reusable, analysis-ready datasets supporting weekly and monthly reporting workflows, eliminating 4–6 hours of manual data wrangling per stakeholder cycle.

PROJECTS

Data Quality-Aware Agent Governance | Python, Microsoft AGT, PostgreSQL, dbt

- Built a two-layer governance prototype using Microsoft AGT that combines agent authorization with dataset trust checks, blocking actions when target data fails freshness or quality thresholds.
- Contributed to Microsoft AGT through three merged PRs: #1818 added the project to ADOPTERS.md, #2038 fixed Agent SRE tutorial issues after end-to-end Windows verification, and #2114 fixed AgentMesh quickstart scaffold commands, generated code, and dependency instructions.

AI Data Governance & Quality Monitoring System | PostgreSQL, Python, dbt, Power BI

- Engineered a governance layer monitoring 150K+ daily events across ingestion pipelines, enforcing schema validation and systematized cross-source reconciliation for 90% of data sources.
- Deployed real-time anomaly detection cutting incident response time by 25%; built AI-powered Power BI dashboard translating data failures into actionable explanations for non-technical stakeholders.

Decision Intelligence Experimentation Platform | PostgreSQL, Python, Streamlit

- Architected an end-to-end A/B testing platform simulating 50K+ users and 280K+ events; designed hypothesis testing pipelines (z-test, t-test, Mann-Whitney) surfacing a conversion-vs-revenue trade-off that informed a go/no-go product launch decision.
- Designed dbt-powered metrics layer and Streamlit self-serve dashboard, reducing analyst involvement in experiment readouts by ~60%.

Metric Decomposition Engine | PostgreSQL, Python, Pandas, Streamlit

- Developed an automated root-cause analysis system decomposing event-based metrics across dimensions on 1.3M+ GitHub Archive events, ranking each segment's contribution to total metric movement.
- Generated plain-English incident reports via Jinja2 and an interactive Streamlit dashboard, turning multi-hour manual metric investigations into a self-serve workflow running in seconds.

EDUCATION

Masters in Data Science

Jan 2025 – Dec 2026

The University of Texas at Arlington, TX | GPA: 3.8

Bachelor of Engineering in Artificial Intelligence and Machine Learning

Aug 2020 – May 2024

Savitribai Phule Pune University, India | GPA: 8.2