

SOMESH ZANWAR

Arlington, TX | (682) 403-0698 | someshzanwar345@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

M.S. Data Science candidate graduating December 2026 with professional experience in data analytics and hands-on work across AI/ML, analytics engineering, automation, API-backed applications, and agent governance. Experienced in Python, SQL, data validation, reporting automation, machine-learning evaluation, experimentation, and building production-minded systems with testing, controls, auditability, and technical documentation.

EDUCATION

The University of Texas at Arlington - M.S. Data Science | GPA: 3.67/4.00

Jan 2025 - Dec 2026

Savitribai Phule Pune University - B.E. Artificial Intelligence and Machine Learning | GPA: 8.2/10

Aug 2020 - May 2024

SKILLS

Programming & Data: Python, SQL, pandas, NumPy, PostgreSQL, dbt, Excel, data transformation, validation, feature engineering

AI & Analytics: scikit-learn, XGBoost, classification, cross-validation, model evaluation, A/B testing, hypothesis testing, SHAP/LIME

Engineering: FastAPI, SQLAlchemy, REST/JSON APIs, Docker, Git/GitHub, CI/CD, pytest, Streamlit

BI & Delivery: Power BI, Tableau, KPI analysis, requirements gathering, stakeholder reporting, technical documentation, Agile/Scrum

EXPERIENCE

Junior Data Analyst | IVS Software Solutions | Pune, India

Sep 2022 - Sep 2024

- Analyzed 80,000-150,000 operational records per cycle using SQL and Python; improved reporting accuracy by 30% through validation, reconciliation, and discrepancy resolution.
- Automated recurring reporting workflows, saving 4-6 hours per cycle; optimized SQL joins and incremental loads to reduce data-preparation time by 40%.
- Collaborated with three internal teams on Power BI dashboards, KPI definitions, reusable datasets, and recurring reports, translating reporting needs into reusable analytical outputs.
- Documented metric definitions, transformation logic, validation procedures, and discrepancy resolutions to improve consistency and maintainability of recurring reporting.
- Supported operational data preparation, reconciliation, and reporting across SQL, Python, Excel, and Power BI, strengthening repeatable data-quality and business-reporting workflows.

PROJECTS

UTA OIT Intelligent IT Ticket Routing | Data Science Capstone

Sep 2026 - Present

- Developing a data-driven workflow with UTA's Office of Information Technology to analyze incoming support tickets and improve routing to appropriate service teams.
- Applying Python-based data preparation, text/feature analysis, model evaluation, and reproducible engineering practices within a real operational support workflow.
- Translating stakeholder needs into a documented analytical workflow with emphasis on evaluation quality, maintainability, and practical integration into existing support processes.
- Structuring the work so data preparation, evaluation, and routing logic can be reviewed independently, supporting clearer iteration with technical and operational stakeholders.

MetricTrust Control Plane | Python, FastAPI, SQLAlchemy, pytest

Aug 2026 - Aug 2026

- Built an evidence-bound execution control layer for autonomous analytics agents, enforcing policy at the action boundary rather than relying on model judgment; used signed permits and explicit authorization evidence for consequential actions.
- Implemented 21 automated tests covering policy logic, permit validation, and HTTP behavior; documented deployment, API behavior, failure handling, and troubleshooting for maintainability.
- Designed reviewable failure paths and evidence validation so unauthorized or insufficiently supported actions are blocked before execution, strengthening auditability and operational control.
- Separated model-generated analysis from execution authorization, demonstrating a production-minded pattern for preventing plausible but unsupported agent actions from reaching downstream systems.

DataTrust OS | Python, PostgreSQL, dbt, Streamlit

May 2026 - Jun 2026

- Built an operational data-trust system with 12 automated quality checks, incident lifecycle tracking, downstream lineage, automated CI validation, and trust reporting.
- Connected data-quality failures to downstream impact and supporting evidence, turning failed checks into reviewable incidents rather than allowing invalid data to propagate silently.
- Combined PostgreSQL, dbt, Python, and Streamlit into a reproducible workflow for validation, incident visibility, and decision-oriented trust reporting.
- Designed the workflow around explicit checks and traceable incident states so data-quality problems remain visible from detection through downstream-impact review.

Decision Intelligence Experimentation Platform | PostgreSQL, dbt, Python

Mar 2026 - May 2026

- Built an experimentation workflow over 50,000 simulated users and 280,000+ events, calculating conversion lift, p-values, confidence intervals, and segmented results.
- Modeled and transformed experiment data in PostgreSQL/dbt to produce repeatable analytical outputs rather than one-off analysis.
- Translated statistical evidence into rollout recommendations, connecting experiment results and segment behavior to practical product decisions.
- Used segmented experiment results to distinguish aggregate lift from subgroup behavior, improving the evidence available for rollout and follow-up analysis.

OPEN SOURCE

- Microsoft Agent Governance Toolkit - 5 merged PRs spanning policy/data-quality integration, a dbt evidence adapter, AgentMesh scaffolding, and tutorial/documentation fixes.
- AgentTrust / A2A CTEF - completed contributions across verifier enforcement, RAG scan states, runtime support, identity documentation, and agentic-security evidence; proposed a data-quality trust scheme with three conformance vectors.