Somesh Zanwar

Data Analyst | Analytics Engineer | Graduate Data Science Student
Email: someshzanwar345@gmail.com | Phone: (682) 403-0698 | LinkedIn | GitHub | Portfolio

SUMMARY

Data Science Master's student specializing in **SQL-based analytics, KPI development, and business intelligence dashboards**. Applied **Python (Pandas, NumPy) and SQL** to process large operational datasets, perform **EDA**, and build **analysis-ready data models for reporting and forecasting**. Built automated dashboards and analytical pipelines that **reduce manual reporting and highlight key business trends and performance drivers**.

EDUCATION

Masters in Data Science **Jan 2025 – Dec 2026**
The University of Texas at Arlington, TX | **GPA: 3.8**
Relevant Coursework: Data Mining, Neural Network and Deep Learning, Convex Optimization, Cloud Computing and Big Data, Database Systems, Introduction to Probability and Statistics

Bachelor of Engineering in Artificial Intelligence and Machine Learning **Aug 2020 – May 2024**
Savitribai Phule Pune University, India | **GPA: 8.2**
Relevant Coursework: Artificial Intelligence, Machine Learning, Natural Language Processing, Database Management Systems, Big

TECHNICAL SKILLS

Programming: Python, SQL, R
Databases & Data Warehousing: PostgreSQL, MySQL, BigQuery
SQL & Data Manipulation: Joins, Aggregations, Window Functions, Query Optimization
Data Analysis: Pandas, NumPy, Exploratory Data Analysis (EDA), Data Cleaning, Data Validation
Statistical Analysis: Hypothesis Testing, A/B Testing, Descriptive Statistics
Data Visualization & Business Intelligence: Power BI, Tableau, Dashboard Development, KPI Reporting, Data Storytelling
Data Engineering: ETL Pipelines, Data Transformation, Data Preparation, Data Modeling
Machine Learning & AI: Scikit-learn, XGBoost, Random Forest, K-Means, Feature Engineering, NLP, LLM, RAG
Cloud Platforms: Google Cloud Platform
Tools: Git, Docker, Jira, Salesforce (CRM), Microsoft Office Suite(Excel, PowerPoint, Word)

WORK EXPERIENCE

Junior Data Analyst **Jun 2023 – Sep 2024**
IVS Software Solutions — Pune, India

- Processed and transformed **80K–150K record operational datasets** using **Python (Pandas, NumPy) and SQL** to create analysis-ready data for analytics and reporting.
- Conducted **exploratory data analysis (EDA)** across multi-variable datasets to identify trends, anomalies, and data quality issues, **reducing data processing time by 30%**.
- Wrote and optimized **complex SQL queries (joins, aggregations, filtering)** into streamlined query pipelines, **reducing data preparation time by 40%**.
- Designed **internal dashboards using Streamlit, Matplotlib, and Plotly** to monitor data distributions, anomaly patterns, and operational performance metrics.

PROJECTS

AI Data Governance Model | PostgreSQL, dbt, Python, Power BI **Mar 2026**

- Built an **AI-assisted data governance platform** using PostgreSQL, dbt, and Python to monitor GitHub Archive datasets and enforce automated data quality rules.
- Designed a **governance metadata layer** (dataset registry, rule catalog, rule runs, health scoring, incident tracking) to operationalize dataset monitoring across analytics tables.
- Implemented a **rule execution engine** and **AI-based incident explanation assistant**, enabling automated detection, scoring, and explanation of data quality failures surfaced in a Power BI governance dashboard.

SQL Retail Analytics & Inventory Optimization Dashboard | SQL, Python, Streamlit, Pandas, Plotly **Jan 2026**

- Built a **SQL + Python data pipeline** transforming **541K+ retail transactions** into analysis-ready fact/dimension tables; optimized aggregations and indexing to **reduce dashboard load time by 40%**.
- Developed an automated **Streamlit + Plotly BI dashboard** with KPIs and ABC inventory analysis, identifying **top 20% of products driving 75% of revenue** and eliminating **3–4 hours/week of manual reporting**.

Customer Churn Prediction with SHAP Explainability | Python, Scikit-learn, XGBoost, SHAP, Pandas **Dec 2025**

- Developed an **end-to-end churn prediction pipeline** using **XGBoost on 7,043 customers / 21 features**, achieving **ROC-AUC 0.84 and F1-score 0.66** after feature engineering and class-imbalance handling.
- Applied **SHAP explainability** to interpret model predictions and identify key churn drivers, enabling actionable insights for **high-risk customer retention strategies**.